

Predictive Analytics World / Deep Learning World

Exercises - TensorFlow

1. Locate Fischer's Iris Dataset and prepare it for use by a TensorFlow (TF) program. Encode setosa as (1,0,0), versicolor as (0,1,0), virginica as (0,0,1). Split the dataset into a training file (120 items, first 40 of each species) and a test file (the remaining 10 of each species).

2. Write a program to create a simple 4-5-3 neural network classifier prediction model. Use tanh for hidden layer activation. Set up stochastic gradient descent training. Run your program.

3. Modify your program to a 4-(5-5)-3 architecture.

4. Which statement is most accurate?

- a.) A `tf.placeholder` is similar to a variable in most programming languages.
- b.) A `tf.placeholder` is similar to a `const` in most programming languages.
- c.) A `tf.placeholder` is similar to lambda expression in many other programming languages.

5. Which statement is most accurate?

- a.) A `tf.Variable` is used to hold non-numeric data such as `char` and `string`.
- b.) A `tf.Variable` is used to hold integer data such as `tf.int32` or `tf.int64`.
- c.) A `tf.Variable` is used to hold numeric variables that will get values via training.

6. What is the difference between a `tf.Session` and a `tf.InteractiveSession`?

- a.) There is no real difference -- `tf.InteractiveSession` is now deprecated.
- b.) A `tf.Session` is used for a normal program; a `tf.InteractiveSession` is used with IPython.
- c.) A `tf.Session` uses only program data; a `tf.InteractiveSession` can accept user input during run time.

7. Why does TF typically not use softmax activation on output nodes in a neural network classifier?

- a.) The `tf.nn.softmax_cross_entropy_with_logits_v2()` function applies softmax during training.
- b.) The `tf.nn.cross_entropy()` function applies softmax only if preliminary outputs do not sum to 1.0.
- c.) There is no built-in `softmax()` function in TF.

8. Which statement is most accurate?

- a.) When training using the `tf.Session.run()` function, you should randomize the order of training items.
- b.) When training using the `tf.Session.run()` function, you not should randomize the order of training items.
- c.) When training using `tf.Session.run()` it doesn't make any difference if you randomize training order.

```

# iris_tf.py

import numpy as np
import tensorflow as tf
import os
os.environ['TF_CPP_MIN_LOG_LEVEL']='2'

def main():
    np.random.seed(1)
    tfv = tf.__version__
    print("\nUsing TensorFlow version " + str(tfv))

    print("\nLoading iris train and test data \n")
    train_file = ".\\Data\\iris_train_data.txt"
    test_file = ".\\Data\\iris_test_data.txt"

    train_x = np.loadtxt(train_file, usecols=[0,1,2,3], delimiter=",", skiprows=0, dtype=np.float32)
    train_y = np.loadtxt(train_file, usecols=[4,5,6], delimiter=",", skiprows=0, dtype=np.float32)

    test_x = np.loadtxt(test_file, usecols=[0,1,2,3], delimiter=",", skiprows=0, dtype=np.float32)
    test_y = np.loadtxt(test_file, usecols=[4,5,6], delimiter=",", skiprows=0, dtype=np.float32)

    input_dim = 4; hidden_dim = 5; output_dim = 3

    X = tf.placeholder(tf.float32, shape=[None, input_dim])
    y = tf.placeholder(tf.float32, shape=[None, output_dim])

    ih_wts = tf.Variable(tf.random_uniform([input_dim, hidden_dim], dtype=tf.float32, seed=1))
    h_biases = tf.Variable(0.0, dtype=tf.float32)
    ho_wts = tf.Variable(tf.random_uniform([hidden_dim, output_dim], dtype=tf.float32, seed=1))
    o_biases = tf.Variable(0.0, dtype=tf.float32)

    h_nodes = tf.add(tf.matmul(X, ih_wts), h_biases)
    h_nodes = tf.nn.tanh(h_nodes)
    o_nodes = tf.add(tf.matmul(h_nodes, ho_wts), o_biases)
    y_predict = tf.argmax(o_nodes, axis=1) # 0, 1, 2

    learn_rate = 0.01
    max_epochs = 200
    cee = tf.nn.softmax_cross_entropy_with_logits_v2(labels=y, logits=o_nodes)
    cost = tf.reduce_mean(cee)
    optimizer = tf.train.GradientDescentOptimizer(learn_rate)
    trainer = optimizer.minimize(cost)

    init = tf.global_variables_initializer()
    sess = tf.Session()
    sess.run(init)
    print("Starting training")
    for epoch in range(max_epochs):
        indices = np.arange(len(train_x))
        # np.random.shuffle(indices) # hmm --does this matter?
        for ii in range(len(indices)):
            i = indices[ii]
            sess.run(trainer, feed_dict={X: train_x[i:i+1], y: train_y[i:i+1]})
        train_acc = np.mean(np.argmax(train_y, axis=1) == sess.run(y_predict, feed_dict={X:train_x, \
y:train_y}))
        if epoch > 0 and epoch % 10 == 0:
            print("epoch = %4d, train accuracy = %.4f " % (epoch, train_acc))
    print("Training complete \n")

    test_acc = np.mean(np.argmax(test_y, axis=1) == sess.run(y_predict, feed_dict={X:test_x, y:test_y}))
    print("Accuracy on test data = %.4f " % test_acc)
    sess.close()

    print("\nEnd demo")

if __name__ == "__main__":
    main()

```